

Evolving Expectations in the Enterprise Cloud Market

How AI, Cost Pressure, and Data Control Are Changing the Cloud Buyer's Mindset in UK Public Sector

Version 1.0

Copyright © 2025, Oracle and/or its affiliates

Confidential – Oracle Internal

Introduction

When public cloud services first reached mainstream acceptance a decade ago, most organisations pursued them for pragmatic reasons: lower capital expenditure, pay-as-you-go flexibility, and the promise of near-instant scalability. Those drivers still matter, yet the centre of gravity has shifted. Advances in generative AI, a sharper board-level focus on cost discipline, and rising regulatory scrutiny of data sovereignty now set a richer and more demanding agenda for cloud adoption. Enterprise buyers weigh access to high-end GPU clusters against the predictability of their monthly invoices; they balance global reach with the need to keep sensitive data within jurisdictional boundaries.

For OCI architects, understanding this evolution is essential. The conversations that once revolved around migrating virtual machines now explore model governance, carbon budgets, and multi-cloud resilience. This paper traces how expectations have changed, why they matter for the UK public sector, and how OCI can address them in a clear, evidence-based way.

From Cloud 1.0 to Today

In the first enterprise wave (approximately 2015 – 2020) cloud projects were justified mainly on three pillars: faster delivery of IT services, lower up-front capital spends and simplified standardisation of infrastructure. IDC's CloudView survey captured the sentiment succinctly; "improving agility and security, as well as standardising IT infrastructure" ranked as the top three drivers for adoption, and **87 % of organisations reported at least some hybrid-cloud capability** at that stage. Gartner, Forrester and other analyst houses echoed the theme: public cloud was a tactical means to cut procurement cycles and modernise without rewriting every application.

Enterprises measured success in similarly straightforward terms. McKinsey's analysis of early movers found companies could **reduce IT overhead costs by 30 – 40 %** when workloads were re-platformed rather than simply "lift-and-shifted". Flexera's *State of the Cloud* survey revealed that adopters were already **running 79 % of their total workloads in some form of cloud - 38 % in public clouds, 41 % in private clouds**, while "optimising cloud spend" topped the governance agenda. Typical KPIs of the period therefore read like IT scorecards: percentage of estate migrated, data-centre closures achieved, or headline run-rate savings delivered against baseline infrastructure budgets.

Yet even as those early targets were being met, a more complex set of expectations started to emerge, rooted not just in cost or agility, but in how cloud could unlock new capabilities, manage risk and satisfy board-level scrutiny.

Five Evolving Priorities

What today's buyers put on the table when they evaluate cloud proposals.

1. AI capability

Enterprise demand for generative AI has put unprecedented pressure on GPU supply, driving **53 % year-on-year growth in data-centre capital expenditure during Q1 2025** alone. Buyers now ask three questions up front:

- **Can you give me scale when I need it?** OCI's roadmap includes a *64 000-GPU liquid-cooled NVIDIA GB200* training cluster, announced on Oracle's FY25 Q3 call.
- **Is the full tool-chain ready to use?** Organisations increasingly prefer managed services that bundle frameworks (TensorFlow, PyTorch), vector databases and governance controls rather than assembling components themselves.
- **How do you secure the models and the data?** Isolation boundaries, key-management options and audit trails are now part of every RFP, reflecting lessons from early AI pilots.

2. Cost predictability & FinOps

Cloud "bill shock" has moved up on the board agenda. **Nearly half of enterprises overspent their 2023 cloud budgets, and 59 % expect overruns again in 2024** according to IDC's Cloud Pulse survey. The response is the rapid institutionalisation of FinOps:

- CFOs demand *variance within 5 %* of forecast; few organisations consistently achieve it.
- Reserved-instance commitments, chargeback models and real-time anomaly alerts are now baseline hygiene.
- Some workloads are being repatriated or re-platformed when long-term OPEX exceeds owned-asset TCO, but the dominant trend is optimisation rather than retreat.

3. Data control & sovereignty

Regulation continues to narrow the freedom to store and move data. The **EU Data Act**, in force since 2025, obliges cloud providers to remove "commercial, technical or contractual" barriers to switching and caps exit fees. Although the United Kingdom is no longer an EU member, the Act still touches UK public-sector bodies that supply services or handle data within the EU. In parallel, UK guidance stresses multi-vendor strategies and explicit location control. The combined effect is visible in procurement checklists:

- Preference for platforms that let customers hold their own encryption keys, operate with in-jurisdiction support teams and provide contractual data-residency guarantees.
- Growing interest in dedicated or sovereign cloud regions managed under local legal authority.
- Explicit questions on how data can be exported in open formats without punitive costs.

4. Resilience & continuity

Critical services, from health to transport, now require *multi-region* or even *multi-cloud* fail-over designs. The **UK Government Resilience Action Plan 2025** calls for higher readiness across public infrastructure, explicitly placing resilience obligations on suppliers and operators. Standard expectations include:

- 99.9 % or better uptime SLAs with documented disaster-recovery run-books.
- Geo-redundant deployments that can withstand single-provider outages or geopolitical disruption.
- Continuous testing (chaos days, game-days) to validate fail-over procedures.

5. Sustainability & ESG

Cloud purchasing teams increasingly score proposals on environmental impact as well as price-performance. Hyperscalers have responded with time-bound public targets and workload-level reporting tools:

- **Amazon (AWS)** – matched **100 % of its global electricity consumption with renewables in 2023**, seven years ahead of its 2030 goal.
- **Microsoft (Azure)** – pledges to be **carbon-negative by 2030** and to remove its historic emissions by 2050; publishes an annual environmental-progress report.
- **Google Cloud** – aims to operate on **24 / 7 carbon-free energy in every region by 2030**, exceeding the industry norm of annual matching.
- **Oracle (OCI)** – targets **100 % renewable energy for all cloud regions by 2025** and achieved **86 % renewable coverage across OCI in 2023**; the company is committed to **net-zero emissions by 2050** with a 50 % absolute-emissions reduction by 2030.

As ESG officers join cloud-governance boards, customers now request carbon dashboards, workload-specific emissions factors and evidence that providers influence upstream supply-chain impacts. For us, Oracle's 2025 renewable-energy milestone and net-zero pathway are assets to surface early, especially when competing against more publicised sustainability narratives from other hyperscalers.

Regulation and Sovereign Cloud

Public-sector architects in the United Kingdom work within one of the most prescriptive regulatory landscapes in the world. Four instruments shape every cloud decision.

1. **The UK GDPR and the Data Protection Act 2018**
After Brexit the General Data Protection Regulation was transposed into domestic law. The Act reinforces principles of purpose limitation, data-minimisation and explicit consent, and permits exemptions only where national-security or vital public interests apply. Any service that stores personal data must show **where** the data will reside, **who** can access it and **how** it can be exported or deleted on demand.
2. **EU Data Act 2024 / 2025**
Although no longer an EU member, many UK authorities handle data that flows to or from European partners. The Data Act, in force since January 2024 and directly applicable from September 2025, compels providers to remove technical and contractual “switching barriers”. It limits exit fees to direct costs and bans them outright from January 2027. Buyers now expect cloud contracts to spell out **portability** terms, **data-egress** pricing and time-boxed **migration support**.
3. **Cabinet Office and DSIT cloud guidance (February 2025)**
The Government Digital Service, part of the Cabinet Office, published *Multi-region cloud and software-as-a-service* guidance that encourages departments to adopt a **controlled multi-region approach**. It states that workloads may be hosted in overseas regions *if* this improves resilience, capacity or access to innovation, provided legal obligations are met. The policy underlines multi-vendor and cross-border strategies rather than “UK-only” hosting in all cases.
4. **National Cyber Security Centre (NCSC) cloud security framework**
The NCSC's collection on cloud security sets baseline controls for **classification**, **encryption**, **logging** and personnel **vetting**. It is widely used for OFFICIAL and SECRET workloads and is referenced in departmental architectural patterns.

These rules drive a clear procurement preference: cloud providers must offer **guaranteed data residency, transparent portability, and locally administered control planes**. Sovereign cloud options, such as our dual-region UK Government Cloud (OC4) or Dedicated Region deployments, address that need by giving departments contractual assurance that **UK law governs data access** and that **exit paths remain open**.

Multi-Cloud & Strategic Hedging

Enterprises rarely place all workloads with a single provider any more. In the **Flexera 2025 State of the Cloud** recap, **70 % of organisations report a hybrid model using at least two public clouds, and the average enterprise consumes 2.4 public-cloud platforms**. The motivations are practical rather than ideological:

- **Risk distribution and resilience** – critical services can fail over to a second cloud or remain portable if a geopolitical or commercial disruption occurs.
- **Best-of-breed access** – AI teams may favour a cloud that offers a specific GPU generation or a managed foundation-model service not yet available elsewhere.
- **Commercial leverage** – procurement teams benchmark prices across vendors and use future portability to negotiate committed-spend discounts.
- **Data-gravity management** – large data sets anchor applications; keeping compute close to data avoids expensive transfers and latency penalties.

Interoperability in practice

Oracle invests heavily in friction-free links to other hyperscalers, turning strategic hedging from theory into architecture:

Capability	What it is	Why it matters
OCI–Azure Interconnect	Private, dual-provider network link with zero ingress or egress fees on either side	Allows low-latency (<2 ms round-trip) access to Oracle databases from Azure VMs without unpredictable network charges.
Oracle Database@Azure	Fully managed Oracle Autonomous Database and Exadata services co-located inside Azure data centres ; no latency or egress costs between the services and native Azure workloads	Lets UK public-sector teams keep existing Azure tooling while exploiting Oracle’s flagship database and AI features.
OCI standard egress pricing	Flat, market-leading rates and no charge for intra-region or availability-domain traffic , unlike most providers	Simplifies TCO modelling and discourages “bill shock” when traffic grows.

These features help organisations place data where policy demands and move compute to whichever cloud offers the optimal cost-performance profile at a given moment. For ourselves, they translate into design patterns such as:

- **Train in OCI, infer in Azure** – Run large-scale model training on OCI’s high-density clusters and serve low-latency inference from a nearby Azure region, pairing cost-efficient training capacity with end-user proximity.
- **Sovereign store, elastic burst** – Keep regulated datasets in the UK Government Cloud (OC4) and burst peak-period analytics to a commercial OCI region, preserving data sovereignty while accessing newer services and short-term capacity without year-round over-provisioning.

- **Database@Azure, micro-services on OCI** – Migrate the legacy Oracle database into Database@Azure to avoid major application change, then deploy new micro-services on OCI to exploit the zero-egress interconnect and lower network costs.

Multi-cloud is therefore less about splitting workloads evenly and more about **strategic hedging**; placing each component where it yields the best blend of compliance, performance and economics, with exit paths kept deliberately open.

Outcome-Based Success Metrics

As highlighted early programmes judged cloud differently, assessing progress by how many virtual machines migrated and how much capex was avoided. Those scorecards now look naïve. Cloud steering boards in 2025 track **financial, environmental and operational outcomes** that map directly to organisational performance.

Metric	Typical KPI	Why it matters	Current benchmarks
Financial (FinOps)	• Forecast variance $\leq 5\%$ • <i>Cost-per-inference</i> for AI models • Commit-utilisation ratio	Boards demand predictable run-rates and business-unit charge-backs rather than headline savings.	IDC Cloud Pulse shows ~50 % of buyers overspent budgets in 2023 and 59 % expect overruns again in 2024 . FinOps Foundation's 2025 survey reports that only 15 % of practices fully optimise spend today .
Carbon & ESG	• Kg CO₂e per workload-hour • % workloads on >90 % renewable power • Supplier CDP / SBTi scores	Government and listed companies face mandatory Scope 3 disclosure; cloud emissions now show on CFO dashboards.	Oracle (OCI) reached 86 % renewable electricity in 2023 and targets 100 % by 2025, while AWS hit 100 % in 2023 and Google aims for 24 / 7 CFE by 2030. Analysts note every major provider now offers a customer carbon calculator, though studies also warn of inconsistent methodologies.
Operational resilience	• Recovery time objective (RTO) < 30 min • Dual-region or multi-cloud fail-over tested quarterly	Outages, cyber incidents and supply-chain risk place uptime in board risk registers.	UK Cabinet Office guidance (2025) encourages <i>"controlled multi-region deployment to improve continuity"</i> .

- CFOs **fund and scrutinise** cloud spend; FinOps now reports to Finance in 44 % of enterprises, up from 29 % two years ago.
- CISOs set security and resilience thresholds, often with veto power on data-sovereignty choices.
- Chief Sustainability (or ESG) Officers add carbon ceilings and supply-chain criteria; the FinOps 2025 study shows that **only 3 % of practices currently optimise for carbon**, highlighting a gap architects can close.

Here we need to deeply understand that success arguments start with **cost-per-business-transaction, tonnes CO₂e avoided, or minutes of recovery time** – no longer gigabytes migrated. Designing architectures that expose these metrics (e.g. OCI Cost Analysis, Oracle Cloud carbon-footprint guidance and Oracle Full Stack DR) positions Oracle credibly against peers and resonates with the multi-stakeholder governance now in place.

How the other Hyperscalers Are Responding

For all our benefit I wanted to explore how the four largest providers have each moved to shore-up gaps revealed by the new buying priorities. Here is a side-by-side snapshot that helps Oracle architects anticipate competitor claims.

Theme	AWS	Microsoft Azure	Google Cloud
AI infrastructure & services	Trainium 2–based <i>Trn2</i> clusters (up to 4× v1 performance) for Gen-AI training; Amazon Bedrock AgentCore for secure, enterprise AI agents	Blackwell Ultra GPU VMs arrive late 2025; Azure first to preview GB200 at production scale	Ironwood (7th-gen TPU) for inference; Trillium (6th-gen) now GA for training
Cost-management & FinOps tools	Discount-aware optimisation checks surfaced in Trusted Advisor & Cost Optimisation Hub (2025)	Monthly Cost Management releases; Copilot for Azure cost insights; Native FOCUS exports to Microsoft Fabric	FinOps Hub and BigQuery FOCUS export for rules-based cost allocation and AI-driven forecasting
Sovereign & data-control options	Independent AWS European Sovereign Cloud with separate control-plane, EU-resident staff, launch by end 2025	New EU sovereign cloud services announced 2025; UK G-Cloud certs remain but residence guarantees under scrutiny	Expanded EU Sovereign Cloud with local operators and Data-Boundary controls 2025
Resilience & multi-cloud interop	Strong cross-region patterns; Outposts & Local Zones for edge continuity.	Azure Arc stretches governance onto any infrastructure; joint Microsoft-Oracle “Fabric” roadmap.	Anthos and Multi-Cluster Ingress for hybrid control-planes.
Sustainability pledges	Matched 100 % of global electricity with renewables in 2023 , seven years early	Carbon-negative by 2030 goal intact, but 2024 report shows 23 % rise in emissions since 2020 driven by AI demand	Operate on 24/7 carbon-free energy in every region by 2030 ; Europe > 90 % CFE in some regions by 2025

- All providers now publish FinOps-grade cost and carbon dashboards.
- Sovereign cloud is no longer niche, each vendor offers (or plans) an isolated EU solution.
- At extreme AI scale, other than OCI only AWS currently disclose clusters above 100k next-gen GPUs.

Oracle OCI Benchmark Table

The table below reveals how strongly Oracle stands on each of these aspects:

Theme	Key Facts & Capabilities
AI Infrastructure & Services	• Supercluster scale: supports up to 131 072 NVIDIA Blackwell GPUs (largest public cluster announced). • Next wave: 64 000-GPU liquid-cooled GB200 cluster under construction (first capacity FY26 Q1). • Managed Gen AI: OCI Generative AI (Cohere & Llama 2) GA Oct 2024; integrates with Autonomous DB vector search & OCI Data Science. • Network fabric: RDMA v3 (3.6 μs, 3200 Gb/s) – MLPerf GPT-3 175 B benchmark scored 1.9 $\times$ Azure (Dec 2024).
Cost-Management & FinOps	• FinOps Hub, Cloud Advisor and 6-hourly FOCUS cost reports built-in; “what-if” RI planner keeps forecast variance < 5 %. • Automatic anomaly detection (≥ 10 % spike alert inside 6 h). • Zero egress fees on OCI–Azure Interconnect; flat global egress tiers (first 10 TB \$0.0085/GB). • Premier member, FinOps Foundation.
Sovereign & Data-Control Options	• Dual-region UK Government & Defence Cloud (London + Newport) for OFFICIAL-SENSITIVE workloads, UK-staff-only ops. • Dedicated Region / Oracle Alloy – full OCI stack on-prem or partner DC, customer-controlled keys & personnel. • BYOK/HYOK & external KMS for all storage; NCSC CS-14 compliant. • EU Sovereign Cloud roadmap: Frankfurt & Madrid regions in H1 2026.
Resilience & Multi-Cloud Interop	• OCI–Azure Interconnect (< 2 ms, zero egress/ingress); Database@Azure places Exadata inside Azure DCs. • Dual-network Clos + remote-direct-storage replication between UK Gov regions; always-free outbound traffic within the pair. • Documented DR fail-over between London and Newport; RTO objectives available under NDA. • Cross-cloud Autonomous DB supported on Azure, AWS & OCI; Data Guard across clouds in preview (H2 2025).
Sustainability Pledges	• 100 % renewable energy for all OCI regions by 2025 (86 % achieved globally in 2023). • Net-zero (Scopes 1-3) target 2050; Scope 1 & 2 carbon-neutral by 2025. • 5.4 $\times$ performance-per-watt gain from 2019 $\rightarrow$ 2024 across Gen 1–3 compute shapes. • Member, iMasons Climate Accord; publishes Environmental Product Declarations (EPDs) for latest shapes (X10M, BM.GPU.C25.8, future B200).

Oracle is More Differentiated than any Other

Oracle has taken a deliberately concentrated path, investing in sovereign-ready capacity, removing cost barriers to multi-cloud designs and scaling AI infrastructure through deep ecosystem partnerships. Further evidence can be found in the table below that empowers your conversations and attest to this statement.

Differentiator	What Oracle Has Done	Why It Matters for UK Public Sector
£ 5 billion UK build-out	Committed to spend £ 5 billion (~\$ 5 billion) over five years to expand OCI regions, GPU clusters and renewable-energy projects in Britain.	Signals long-term presence and provides capacity headroom for departments anticipating AI-driven growth.
Founding role in Stargate & OpenAI deal	Oracle named as initial equity funder and technology partner in OpenAI's Stargate initiative; reported \$ 30 billion-per-year contract to supply ≈ 4.5 GW of AI datacentre power to OpenAI.	Demonstrates ability to provision hyperscale AI infrastructure that meets the security and governance bar of one of the world's most scrutinised AI developers.
Strategic alliance with NVIDIA	Building a 64 000-GPU, liquid-cooled NVIDIA GB200 super-cluster and offering > 160 NVIDIA AI Enterprise tools and NIM micro-services natively in OCI.	Architects can propose state-of-the-art GPU capacity without waiting for overseas supply; software stack is pre-integrated and supported.
Zero-egress, zero-latency multi-cloud	OCI-Azure Interconnect waives network charges and delivers sub-2 ms latency; Database@hyperscaler places Exadata inside competitor cloud facilities, still billed at OCI rates.	Lets departments observe Cabinet Office multi-region guidance while avoiding the "data-gravity tax" that inflates TCO on other clouds.
Sovereign deployment choices	Dual-region (London + Newport) UK Government & Defence Cloud (OC4) for OFFICIAL-SENSITIVE workloads; Dedicated Region and Alloy for on-prem or partner-hosted control planes.	Meets localisation and personnel-clearing requirements without sacrificing OCI's full service catalogue.
Flat, transparent network pricing	Uniform low egress rates worldwide; intra-region and AD traffic free, cross-cloud egress waived on Azure Interconnect.	Simplifies FinOps forecasting and aligns with EU Data Act mandates to remove switching penalties.
Renewable-energy target on short fuse	86 % renewable electricity in 2023; 100 % target for all OCI regions by 2025.	Supports departmental net-zero roadmaps and avoids long-distance data transfers to reach green regions.

OCI Architects' Playbook for Enterprise Cloud Conversations

Based on the up-to-date assessments provided in this paper you should be armed with practical prompts to frame value, risk and cost when advising UK-public-sector customers.

1. Start with the outcome scorecard

- **Clarify success measures upfront:** cost-per-inference, carbon ceiling (kg CO₂e / workload-hour), RTO/RPO, citizen-service KPIs.
- **Map to telemetry:** OCI Cost Analysis, 6-hour FOCUS exports, Carbon Footprint report, Built-in Service Level Objectives.
- **Baseline and track:** agree a starting metric and a review cadence (e.g. monthly FinOps & ESG checkpoint).

2. Quantify AI scale and spend discipline

- **Quote tangible capacity:** GB200 cluster size, Blackwell GPU count, total GW power envelope; reference the OpenAI and NVIDIA deals as proof of future headroom.
- **Show cost controls in action:** zero-egress interconnect, flat egress tiers, anomaly alerts, RI “what-if” planner.
- **Connect to FinOps practice:** demonstrate how headroom and flat tariffs reduce forecast variance versus burst pricing on peer clouds.

3. Address lock-in and multi-cloud upfront

- **Demonstrate live paths:** OCI–Azure Interconnect, Database@hyperscaler, bidirectional DR between OCI and Azure with no network surcharge.
- **Produce an exit plan:** scripted export in open formats, OCI Object Storage lifecycle policies, Terraform modules for redeploy elsewhere; satisfy UK procurement and EU Data Act portability checklists.
- **Reference tooling:** OCI Resource Manager (Terraform), Cloud Guard, Data Transfer Service for bulk extraction.

4. Match sovereignty to classification

- **Confirm data class first:** OFFICIAL, OFFICIAL-SENSITIVE, SECRET.
- **Default to OC4 where feasible:** every UK Public Sector customer is OC4-eligible; dual-region resilience, UK-staff ops centre.
- **Escalate only when needed:** propose Dedicated Region or Alloy for personnel-cleared, air-gapped mandates; recommend multi-region public OCI for widest service catalogue and rapid GPU availability.
- **Surface controls:** BYOK/HYOK, external KMS, Data Safe, Audit logs exported back to OC4 even when bursting.

5. Tie sustainability to risk management

- **Quote firm milestones:** 100 % renewable energy for all OCI regions by 2025 (86 % reached 2023); 5.4 × perf-per-watt improvement 2019-24.
- **Compare objectively:** show carbon-intensity tables for OCI London versus peer regions at AWS, Azure and Google; reference published EPDs and iMasons Climate Accord membership.
- **Embed in governance:** add carbon forecast to the same dashboard that shows cost and uptime.

6. Prove operational resilience

- **Fail-over design:** Dual-region UK Government Cloud offers documented disaster-recovery fail-over between London and Newport; RTO objectives are available under NDA.
- **Multi-cloud DR option:** Cross-cloud Autonomous Data Guard preview is available for Oracle ↔ Azure or Oracle ↔ AWS protection.
- **Observability stack:** Logging Analytics, Service Connector Hub, and events pushed to the department's SIEM.

7. Engage every stakeholder early

- **Workshop agenda:** Day 0 alignment with CFO/FinOps, CISO, ESG lead, programme owner.
- **Shared-responsibility matrix:** Oracle covers hardware security, green-energy procurement, network isolation; customer covers data lifecycle, application resilience, policy compliance.
- **Decision record:** keep a one-page log that captures chosen region, portability plan, carbon target and cost guardrails, referenced at every steering meeting.

Applied consistently, this playbook translates the benchmark facts: sovereign control, flat pricing, world-scale GPU clusters, into choice-ready architectures that satisfy UK public-sector governance and achieve measurable business outcomes.

Conclusion

Enterprise cloud decisions now hinge on five intertwined priorities: AI scale, cost predictability, data sovereignty, operational resilience and sustainability. The research shows that Oracle has answered each priority with demonstrable assets including Blackwell-class Superclusters, zero-egress multi-cloud links, dual-region UK Government Cloud, built-in FinOps tooling and a near-term 100% renewable-energy goal. The Oracle OCI Benchmark Table distils those facts into a single reference for architects, replacing assumption with evidence.

Advisory conversations must still translate features into outcomes. The OCI Architects' Playbook provides that bridge, outlining stakeholder questions, design trade-offs and compliance checkpoints that align with UK public-sector policy. When architects pair this playbook with the benchmark data they can frame OCI not as an all-or-nothing proposition, but as an adaptable component within a multi-cloud, regulation-conscious strategy.

Continuous learning remains essential. GPU roadmaps, FinOps practices and sovereignty rules evolve faster than project lifecycles. The next paper in this series, *Sovereignty and Trust in Cloud*, will examine the legal and operational foundations that underpin those rules. Mastery of that material will round out the toolkit needed to guide customers through an era where technology choice and jurisdictional trust are inseparable.

Copyright © 2025, Oracle and/or its affiliates. This document is provided for information purposes only, and the contents hereof are subject to change without notice. This document is not warranted to be error-free, nor subject to any other warranties or conditions, whether expressed orally or implied in law, including implied warranties and conditions of merchantability or fitness for a particular purpose. We specifically disclaim any liability with respect to this document, and no contractual obligations are formed either directly or indirectly by this document. This document may not be reproduced or transmitted in any form or by any means, electronic or mechanical, for any purpose, without our prior written permission.

Oracle, Java, MySQL, and NetSuite are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.